





Faster model requests do not always produce faster end-to-end episodes



40 timed episodes per model/task; different budgets. Other time includes execution, I/O and scheduling/contention.